

Announcements

Reminder: ps3 due Thursday 10/8 at midnight (Boston)

Assignment Project Exam Help

- ps4 out Thursday, due 10/15 (1 week)
- Lab this week – neural network learning
- ps3 self-grading form out Monday, due 10/19

<https://powcoder.com>

Add WeChat powcoder

A digital illustration of a neural network. It features a central neuron with a purple nucleus and a blue cytoplasm, surrounded by a dense web of blue, glowing axons and dendrites. Several bright orange sparks or light points are scattered throughout the network, suggesting active signal transmission. The background is dark, making the glowing elements stand out.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Neural Networks III

Today: Outline

- **Neural networks cont'd**

Assignment Project Exam Help

- **Types of networks:** Feed-forward networks, convolutional networks, recurrent networks

**<https://powcoder.com>
Add WeChat powcoder**

- **ConvNets:** multiplication vs convolution; filters (or kernels); convolutional layers; 1D and 2D convolution; pooling layers; LeNet, CIFAR10Net



Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

Neural Networks III

Network Architectures

Neural networks: recap




Learn parameters via gradient descent

$$\min_{\Theta} J(\Theta)$$

Assignment Project Exam Help
<https://powcoder.com>

Add WeChat powcoder

Backpropagation efficiently computes cost (forward pass) and gradient (backward pass)




$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$$


Network architectures

Feed-forward


Fully connected



Recurrent



Convolutional



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

Neural Networks III

Convolutional Architectures

Multiplication vs convolution

Input



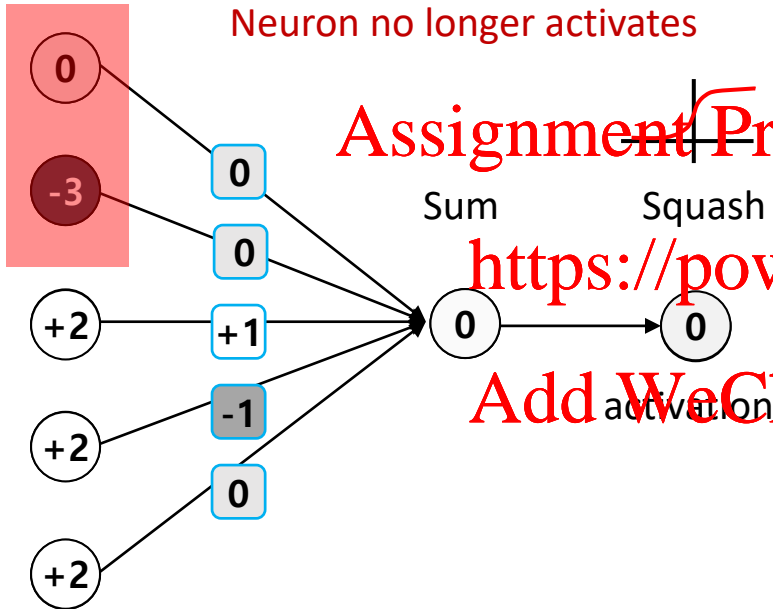
- Recall, a neuron can be thought of as learning to spot certain features in the input
- E.g., this neuron detects change from high to low (light to dark) between 3rd and 4th inputs

<https://powcoder.com>
Add WeChat powcoder

Multiplication vs convolution

Input

What if the change happens
between 1st and 2nd inputs?
Neuron no longer activates



- Must have a new neuron for each new location of pattern???

- This is not efficient

- Solution: use convolution instead of multiplication

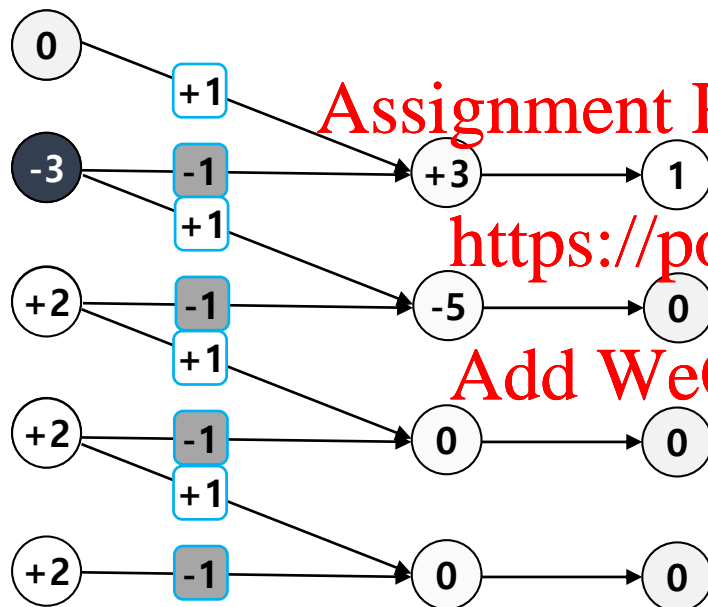
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Multiplication vs convolution

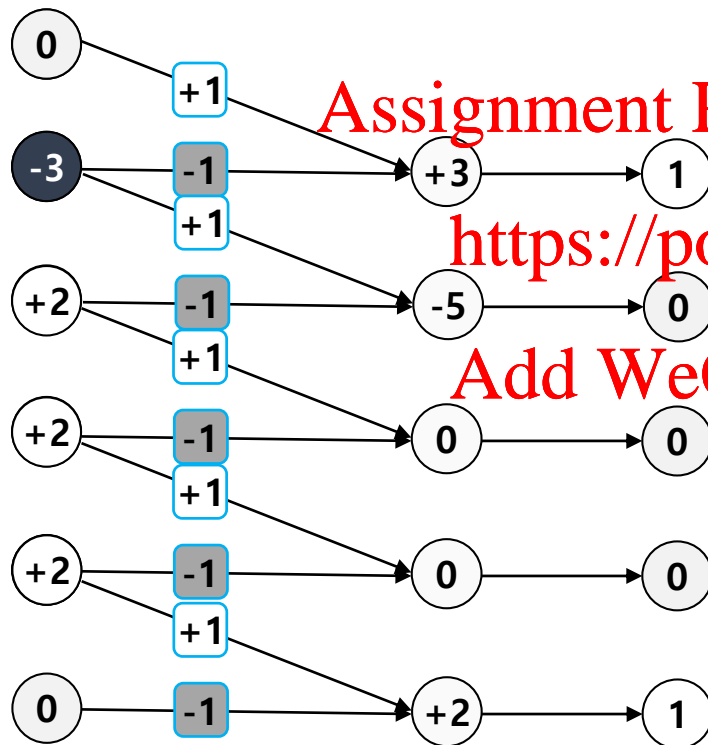
Input



- New weights are of size 2 x 1; called **filter**, or **kernel**
- New output is the size of input minus 1 because of **boundary**
- New convolutional neurons all share the same weights! This is much more efficient; we learn the weights once instead of many times for each position

Multiplication vs convolution

Padded
Input



- New output is the size of input minus 1 because of **boundary**
- We can fix the boundary effect by padding the input with 0 and adding one more neuron

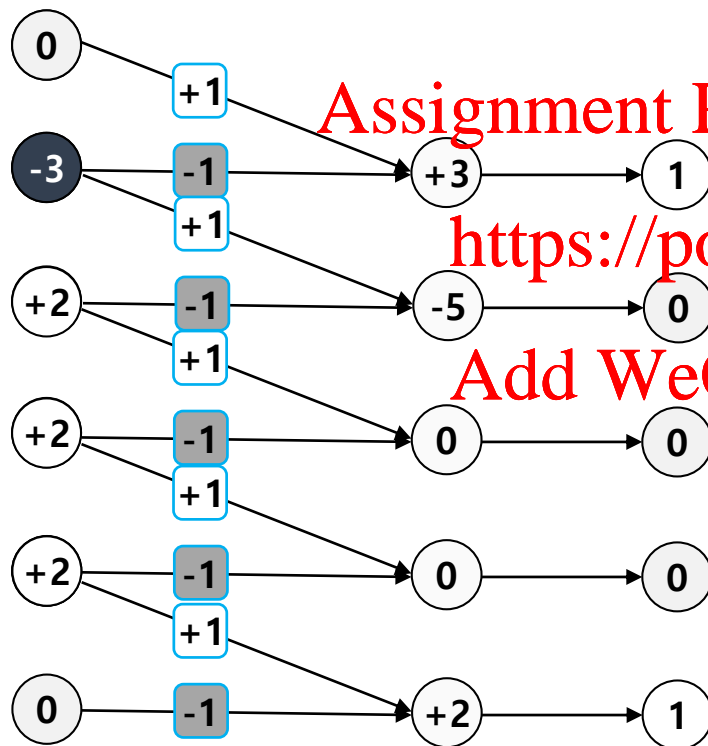
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Multiplication vs convolution

Padded
Input



- Note, we move the filter by 1 each time, this is called **stride**




Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Multiplication vs convolution

Padded
Input



- Note, we move the filter by 1 each time, this is called **stride**
- Stride can be larger, e.g. here is stride 2


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Multiplication vs convolution

Padded
Input




- We can add another filter, this time to detect opposite change with weights $[-1 \ +1]$
- Unique filters are called **channels**

<https://powcoder.com>

Add WeChat powcoder

Multiplication vs convolution



- We can add another filter, this time to detect opposite change with weights $[-1 \ +1]$
- Unique filters are called **channels**


<https://powcoder.com>

Add WeChat powcoder


Multiplication vs convolution

simplified view

Padded
Input



Channels



To summarize, this layer has

- Input 5×1 , padded to 6×1
- Kernel 2×1 with weights $[+1, -1]$
- Stride 2
- Output 3×1
- No. channels K

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Assignment Project Exam Help

<https://powcoder.com>

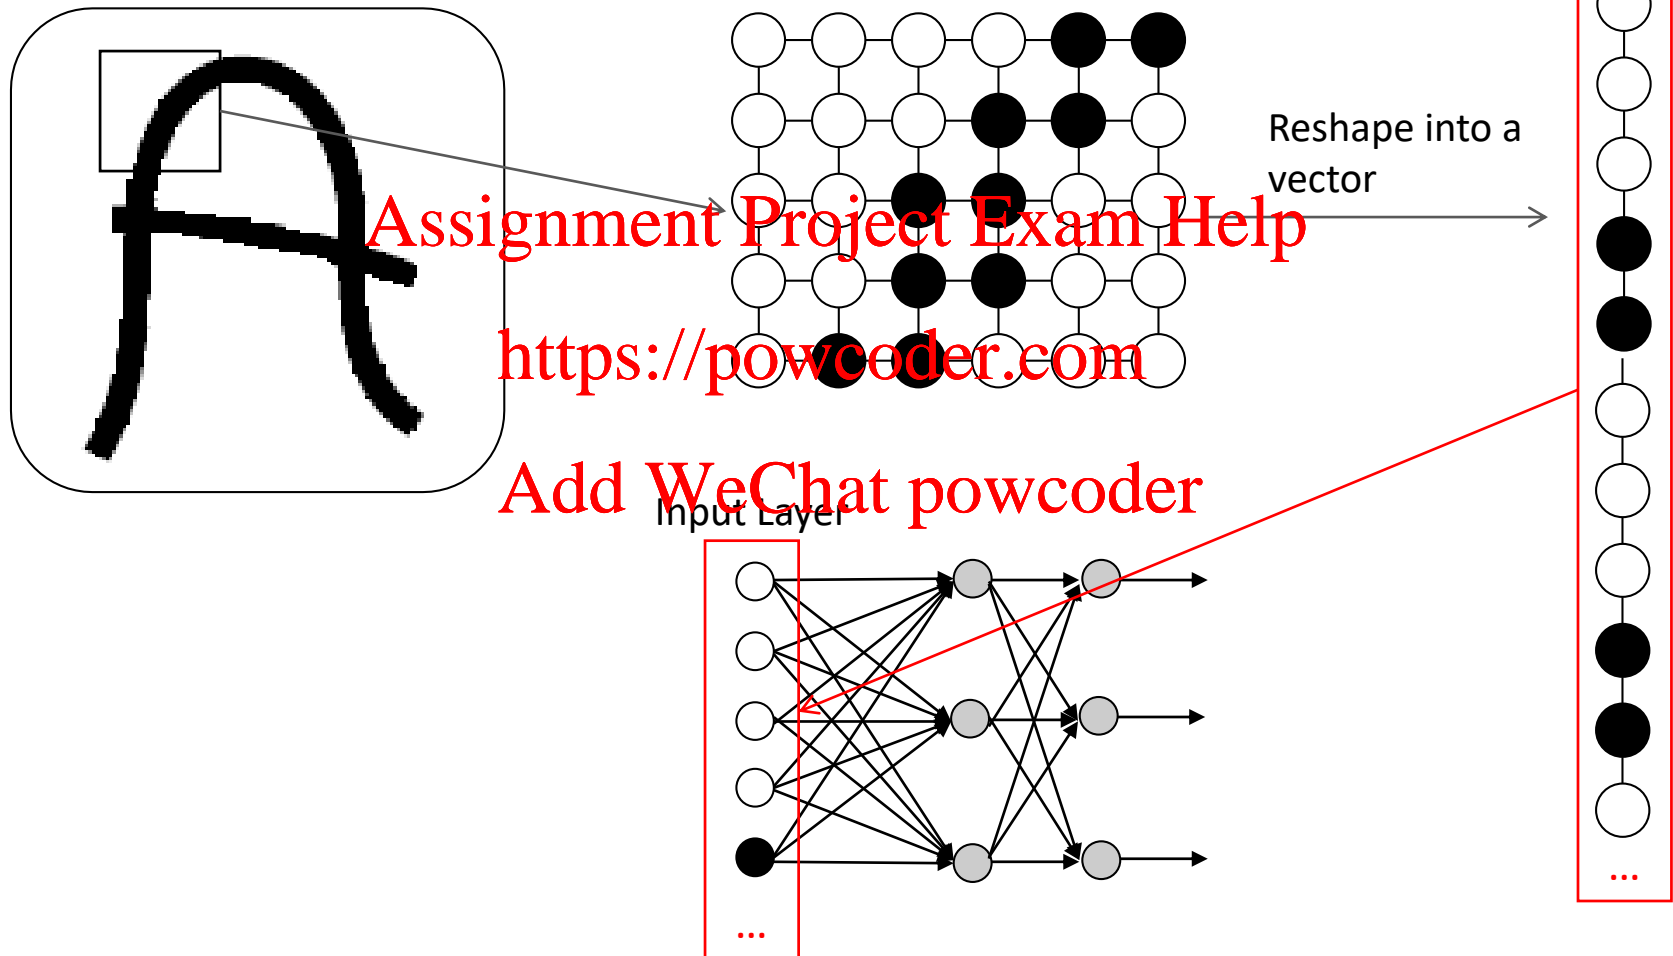
Add WeChat powcoder

Convolutional Neural Networks

For images and other 2-D signals


Representing images

Fully connected



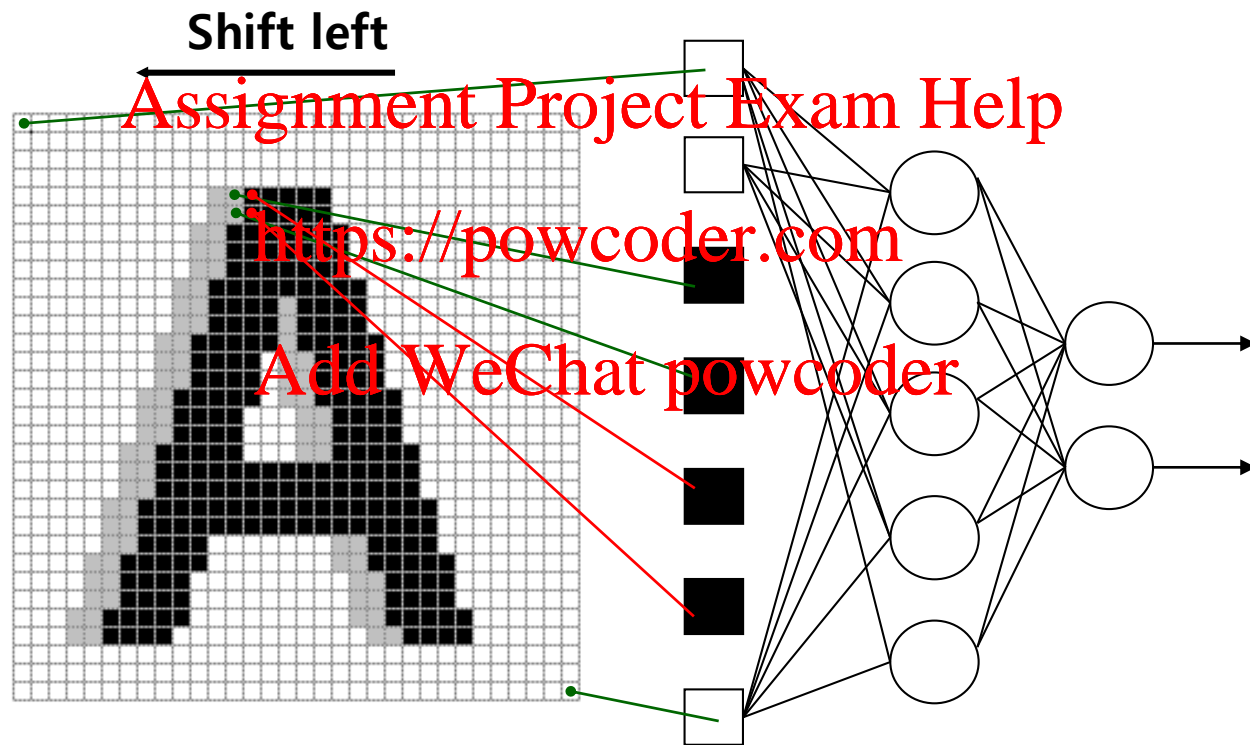
2D Input: fully connected network

Vectorize input by copying rows into a single column



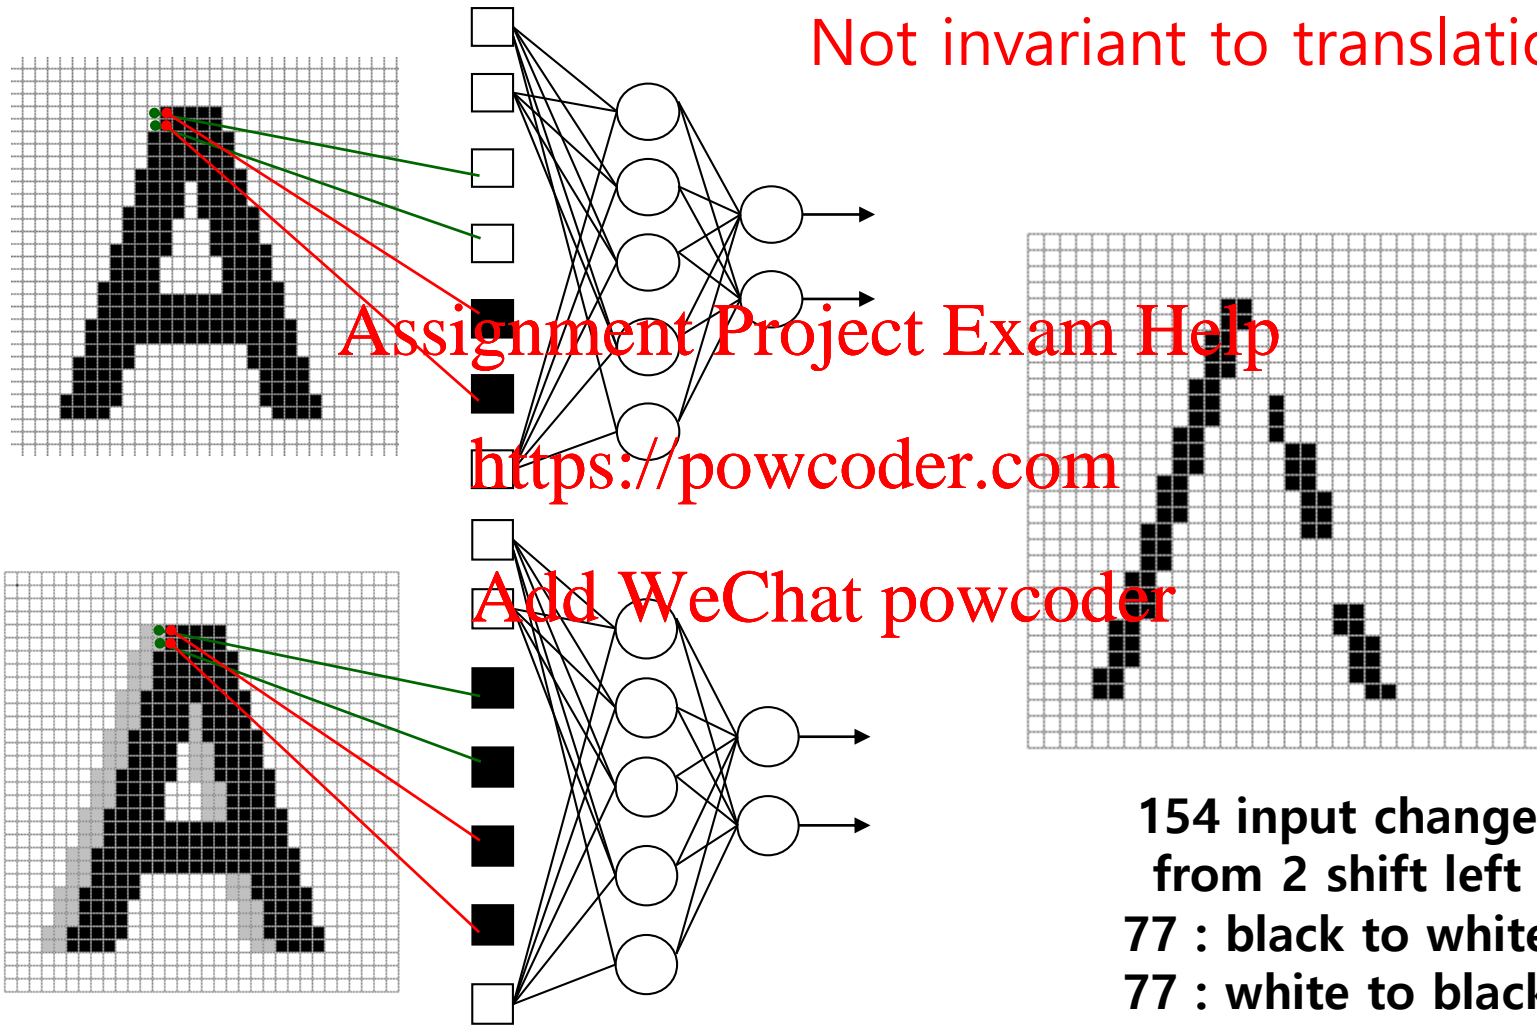
2D Input: fully connected network

Problem: shifting, scaling, and other distortion changes location of features



2D Input: fully connected network

Not invariant to translation!



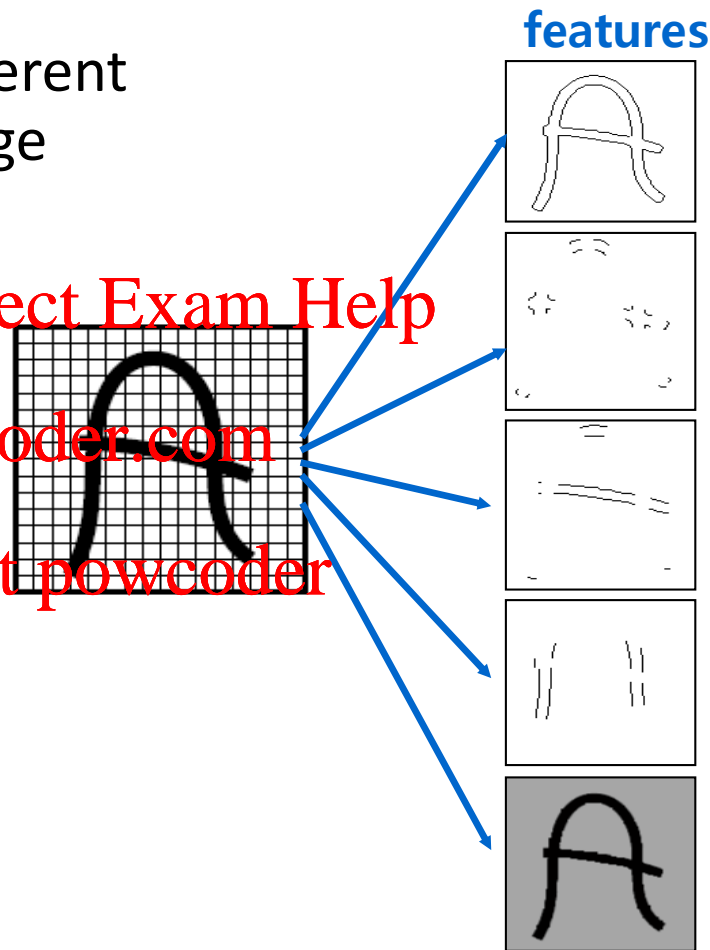
Convolution layer in 2D

- detect the same feature at different positions in the input, e.g. image
- preserve input topology


Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder



Convolution layer in 2D



Input image




Output map

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

Convolution layer in 2D



Input image




Output map

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Output looks like
an image

$$a = f(w_{11}x_{11} + w_{12}x_{12} + w_{13}x_{13} + \cdots w_{33}x_{33})$$

What weights correspond to these output maps?

*These are output maps
before thresholding*

*Hint: filters look like the
input they fire on*





Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


$$\frac{\partial f}{\partial x}(x, y)$$

∂x



$$\frac{\partial f}{\partial y}(x, y)$$

∂y



What will the output map look like?



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



filter

Input

What will the output map look like?

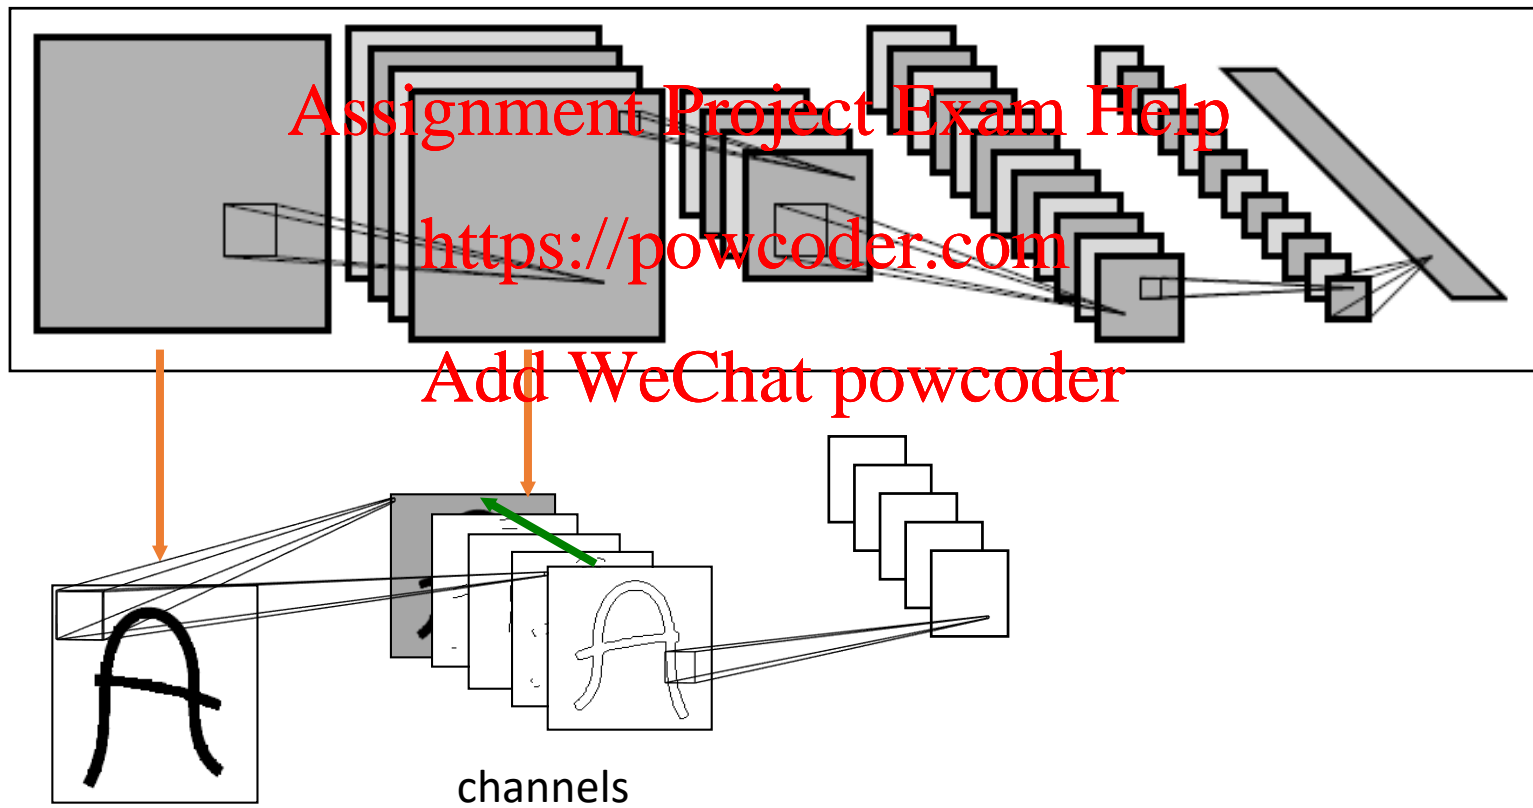


filter

Output

Stacking convolutional layers

- Each layer outputs multi-channel **feature maps** (like images)
- Next layer learns filters on previous layer's feature maps




Pooling layers

- Convolution with stride > 1 reduces the size of the input
- Another way to downsize the feature map is with **pooling**
- A pooling layer subsamples the input in each sub-window
 - **max-pooling**: chose the max in a window
 - **mean-pooling**: take the average

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder




Pooling layer

- the **pooling** layers reduce the spatial resolution of each feature map
- Goal is to get a **certain degree** of **shift** and **distortion**

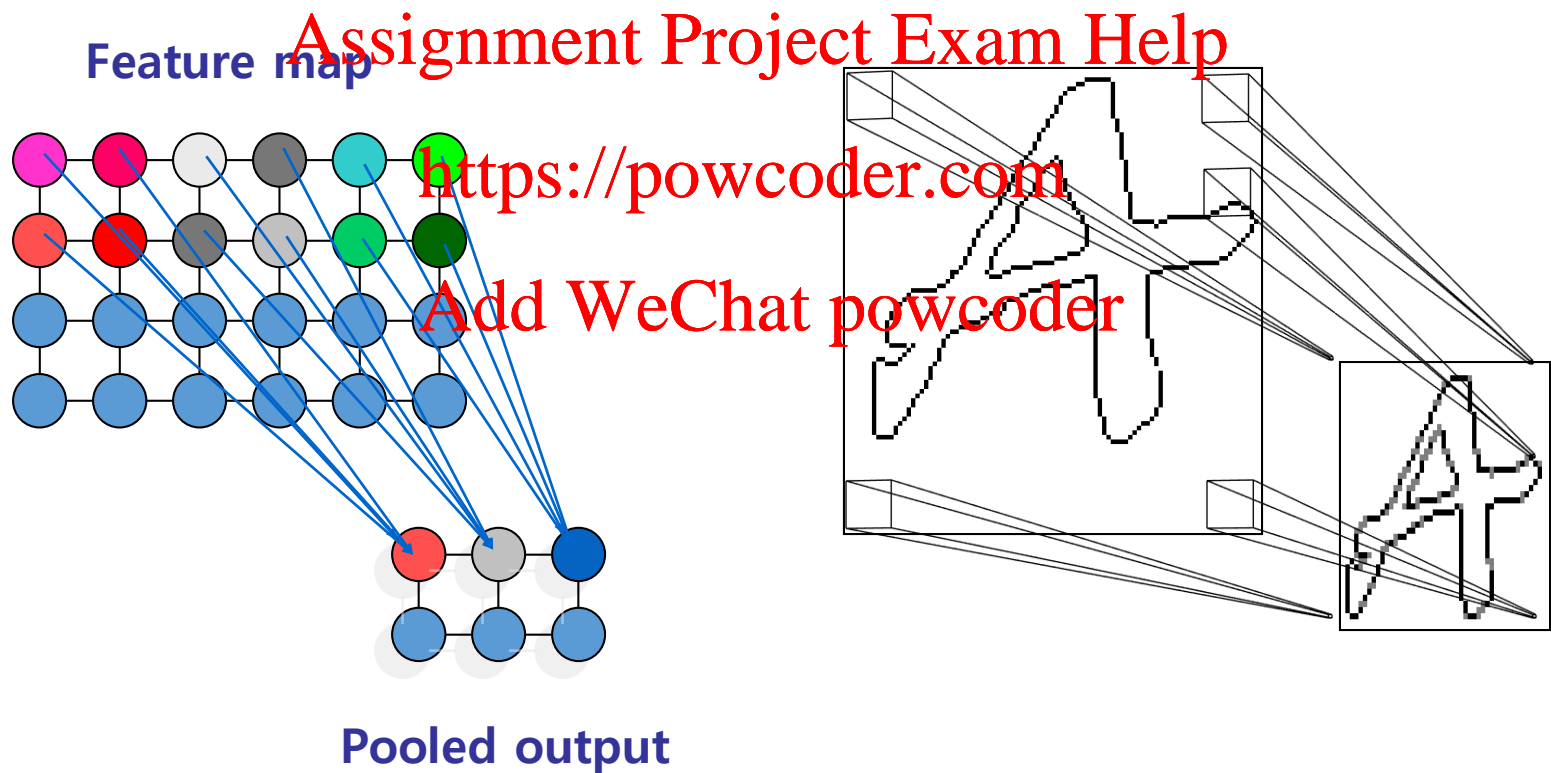


Distortion invariance




Pooling layer

- the **weight sharing** is also applied in pooling layers
- for mean/max pooling, no weights are needed




Putting it all together...



Convolutional Neural Network

A better architecture for 2d signals



LeNet

Deep Convolutional Networks

The Unreasonable Effectiveness of Deep Features



Maximal activations of pool_5 units

[R-CNN]

Rich visual structure of features deep in hierarchy.



conv_5 DeConv visualization

[Zeiler-Fergus]



Assignment Project Exam Help

<https://powcoder.com>




Add WeChat powcoder

Convolutional Neural Nets

Why they rule


Why CNNs rule: Sparsity

- CNNs have sparse interactions, because the kernel is smaller than the input
- E.g. in thousands or millions pixel image, can detect small meaningful features such as edges
- Very efficient computation!
 - For m inputs and n outputs, matrix multiplication requires $O(m \times n)$ runtime (per example)
 - For k connections to each output, need only $O(k \times n)$ runtime
- Deep layers have larger effective inputs, or **receptive fields**



Why CNNs rule: Parameter sharing

- Kernel weights are shared across all locations
- Statistically efficient – learn from more data
- Memory efficient – store only k parameters, since $k \ll m$, this is much smaller than $m \times n$.




*parameter used
multiple times*

*parameter used
only once*

Why CNNs rule: Translation invariance

- Output is invariant to translation of input
 - spatial translation for images
 - temporal translation for time sequences
- useful when some function of a small local window is useful when applied to multiple input locations
- Note, not invariant to other transformations of input, such as large image rotation
- Pooling provides additional invariance to distortions





Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

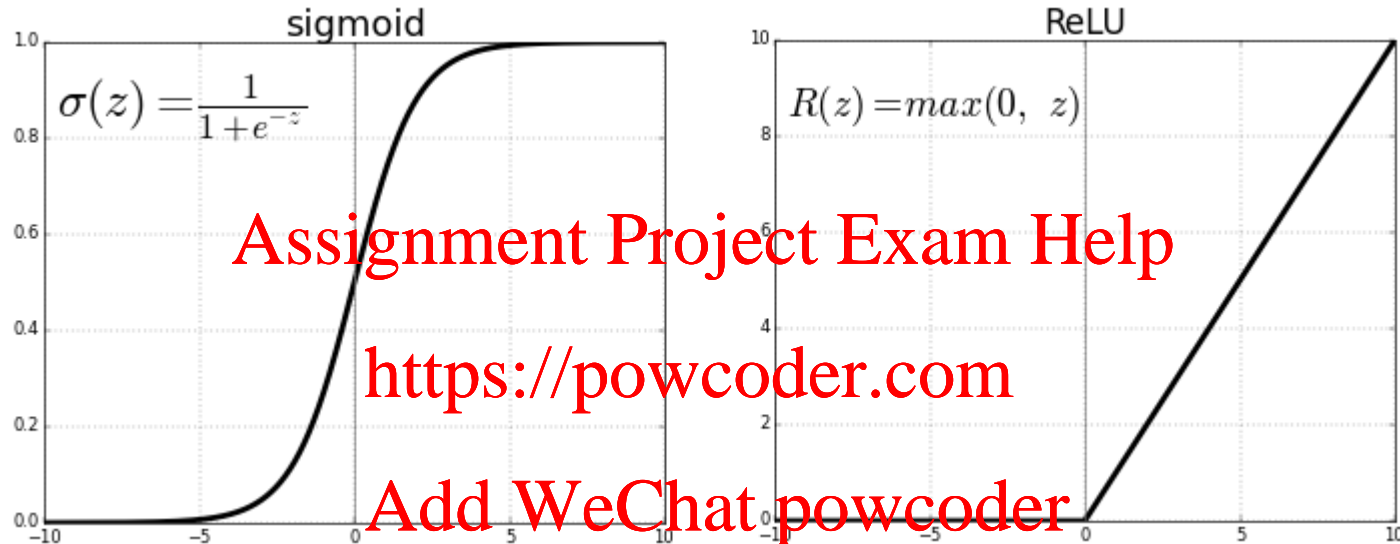
Convolutional Neural Nets

Example


CIFAR-10 Demo ConvJS Network



ReLU: rectified linear unit




ReLU function $g(x) = \max(0, x)$



input (32x32x3)




filter size 5x5x3, stride 1



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder




input (32x32x3)



RELU

conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$


filter size 5x5x3, stride 1



Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder




input (32x32x3)



filter size 5x5x3, stride 1



conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$




pool (16x16x16)
pooling size 2x2, stride 2

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



filter size $5 \times 5 \times 3$, stride 1




input ($32 \times 32 \times 3$)




conv ($32 \times 32 \times 16$) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$

pool ($16 \times 16 \times 16$)
pooling size 2×2 , stride 2



filter size $5 \times 5 \times 16$, stride 1




<https://powcoder.com>
Add WeChat powcoder

• • •


RELU



conv ($16 \times 16 \times 20$) params: $20 \times 5 \times 5 \times 16 + 20 = 8020$



pool ($8 \times 8 \times 20$)
pooling size 2×2 , stride 2



input (32x32x3)



One more conv+RELU+pool:

conv (8x8x20)

filter size 5x5x20, stride 1

relu (8x8x20)

pool (4x4x20)

pooling size 2x2, stride 2

parameters: $20 \times 5 \times 5 \times 20 + 20 = 10020$

Assignment Project Exam Help

<https://powcoder.com>

softmax (1x1x10)

Add WeChat powcoder

dog

car

cat

:

Testing the network

- Show top three most likely classes



<http://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html>

Next Class

Neural Networks IV: Recurrent Nets:

recurrent networks; training strategies

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder